

Bayesian exposure-response modeling for binary data

Outline

- ▶ Brief introduction to Bayesian analysis
- ▶ Fitting models
- ▶ Model checking

Approach to Bayesian modeling in this course

- ▶ For this series of classes we are going to use Stan to do Bayesian modeling
 - ▶ Stan is a probabilistic programming language for fitting Bayesian models.
 - ▶ By default it uses Hamiltonian Monte Carlo (HMC), specifically a variation called the no U-turn sampler (NUTS).
 - ▶ We will go into more details about HMC/NUTS later in the course
- ▶ Start with using `brms` as our gateway to Stan
 - ▶ `brms` is a package that enables simple fitting of many types models that you can fit using `glm`, `survreg`, `lme`, `nlme`, etc.
 - ▶ Allows quick access to Bayesian inference
- ▶ Later we will program our own Stan models and run them with `rstan`
 - ▶ Learn more about the Stan language
 - ▶ Fit models that are not supported in `brms`

A brief review of Bayesian inference

Bayes Rule is the basis for inference about model parameters θ given data y and prior knowledge about model parameters $p(\theta)$:

$$p(\theta | y) = \frac{p(\theta)p(y | \theta)}{p(y)} = \frac{p(\theta)p(y | \theta)}{\int p(\theta)p(y | \theta)d\theta} \\ \propto p(\theta)p(y | \theta)$$

Goals:

- ▶ Inference about θ or a function of θ
- ▶ Predictions of future observations
- ▶ The posterior summarizes what we know about θ , but typically we can't express $p(\theta | y)$ in closed-form.
 - ▶ We'll use Markov Chain Monte Carlo to obtain samples from $p(\theta | y)$.

Bayesian modeling/inference process using MCMC

1. Construct a model for the data, conditional on parameters θ , $p(y | \theta)$
2. Construct a prior distribution for θ , $p(\theta)$
 - ▶ Ideally based on all available evidence/knowledge (or belief)
 - ▶ Or deliberately select a non-informative (or weakly informative) prior
3. Sample from the posterior distribution for θ , $p(\theta | y)$.
 - ▶ Look at convergence and sampler diagnostics
 - ▶ Use for inferences regarding parameter values
4. Sample from the posterior predictive distribution for y_{new} :
$$p(y_{\text{new}} | y) = \int p(y_{\text{new}} | \theta) p(\theta | y) d\theta.$$
 - ▶ Use for inferences regarding future observations
 - ▶ Sample from $p(y | \theta)$ for values of θ from step 3.

Bayesian ingredients for MCMC sampling from a posterior

- ▶ Data
- ▶ Model for the outcome(s) – the likelihood
- ▶ Models for the parameters – the prior distribution
- ▶ MCMC tool – Stan (via brms or 'rstan')

Ingredients for HMC/NUTS

- ▶ A starting point in the parameter space (initial value, one per chain)
- ▶ Number of MCMC samples used to tune the HMC/NUTS algorithm (`warmup`)
 - ▶ This is not exactly the same as the burn-in in other MCMC algorithms
 - ▶ NUTS uses these samples to adaptively tune the sampler
- ▶ Total number of samples to take, including the warm-up (`iter`)
- ▶ Parameters which inform how the sampler should adapt
 - ▶ Defaults are usually good for 'simple' models
 - ▶ Often need to modify for more complex, hierarchical models
 - ▶ Return to these later in the course

Let's re-fit our AE model using brms

```
mod01_glm <- glm(formula = AE01 ~ CAVGSS + BWT + PTTYPE + SEXTXT,  
                 data = aedat,  
                 family = binomial(link = "logit")  
                 )  
  
mod01_stan <- brm(formula = AE01 ~ CAVGSS + BWT + PTTYPE + SEXTXT,  
                 data = aedat,  
                 family = bernoulli(link = "logit"),  
                 warmup = 500,  
                 iter = 2000,  
                 chains = 4,  
                 init = "random",  
                 cores = 2,  
                 seed = 123)
```


What about the prior distributions?

- ▶ By default, `brms` uses flat, non-informative prior distributions for regression coefficients
- ▶ We can specify priors directly through the `prior` argument.
 - ▶ More to come in a few slides

Model summary

```
summary(mod01_stan)
```

```
. Family: bernoulli
. Links: mu = logit
. Formula: AE01 ~ CAVGSS + BWT + PTTYPTYPE + SEXTXT
. Data: aedat (Number of observations: 180)
. Draws: 4 chains, each with iter = 1000; warmup = 500; t
.         total post-warmup draws = 2000
.
. Population-Level Effects:
.           Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk
. Intercept      -7.32      4.47  -16.50      1.00 1.00
. CAVGSS           0.86      0.19   0.52      1.26 1.00
. BWT              0.04      0.06  -0.08      0.16 1.01
. PTTYPTYPEPT1     1.52      0.75   0.17      3.05 1.00
. PTTYPTYPEPT2     0.88      0.97  -1.05      2.74 1.00
. SEXTXTMALE       0.07      0.89  -1.70      1.84 1.01
.
```





Draws were sampled using sampling(NUTS). For each parameter

MCMC convergence diagnostics

- ▶ Traceplots
 - ▶ Plot of sampled values vs iteration
 - ▶ Look for stationarity and good mixing: fuzzy caterpillar
- ▶ \hat{R}
 - ▶ Heuristically: $\frac{\text{total variance of } \theta \text{ (between and within-chain)}}{\text{average within-chain variance of } \theta}$
 - ▶ Target: $\hat{R} < 1.01$ (sometimes, you'll see a rule of $\hat{R} < 1.05$)
 - ▶ Output: Summary and plot (`mcmc_plot(mod01_stan, type='rhat')`)
- ▶ Effective sample sizes
 - ▶ bulk ESS for assessing posterior means, medians, etc
 - ▶ tail ESS for assessing tail percentiles (5th, 95th)
 - ▶ Target: Depends on your goals
 - ▶ Output: Summary

Traceplots

```
bayesplot::mcmc_trace(mod01_stan)
```



Let's see the default priors in our model

```
prior_summary(mod01_stan)
```

```
.               prior      class      coef group resp dpa
.             (flat)         b
.             (flat)         b      BWT
.             (flat)         b    CAVGSS
.             (flat)         b  PTTYPEPT1
.             (flat)         b  PTTYPEPT2
.             (flat)         b  SEXTXTMALE
. student_t(3, 0, 2.5) Intercept
.      source
.      default
. (vectorized)
. (vectorized)
. (vectorized)
. (vectorized)
. (vectorized)
.      default
```

BRMS centers all predictors in the model

Mathematically, the RHS of the model $y \sim x_1 + x_2$ is

$$b_Intercept + b_x1 \cdot x_1 + b_x2 \cdot x_2$$

Or, equivalently,

$$Intercept + b_x1 \cdot (x_1 - \overline{x_1}) + b_x2 \cdot (x_2 - \overline{x_2})$$

where

$$b_Intercept = Intercept - b_x1 \cdot \overline{x_1} - b_x2 \cdot \overline{x_2}$$

This is the parameterization that brms uses. (***Why do you think that is?***)

We need to specify priors for Intercept, b_x1 and b_x2

To change them use the `set_priors` function

A `Normal(mean=0, sd=5)` prior on all covariate effects:

```
priors_mod01 <- set_prior('normal(0,5)', class='b')
```

A `Normal(0,5)` prior on CAVGSS and a `DoubleExponential` prior on the other effects:

```
priors_mod01_de <- set_prior('normal(0,5)', class='b', coef = 'CAVGSS')  
  set_prior('double_exponential(0,1)', class='b')
```

See Stan functions reference for list of all available distributions.

Workbook Bayes01

- ▶ Model fitting in brms
- ▶ Convergence diagnostics

Model diagnostics

- ▶ We can use similar diagnostics as with likelihood based methods, but now using posterior predictive distributions
 - ▶ Quantile residual plots
 - ▶ Posterior predictive checks


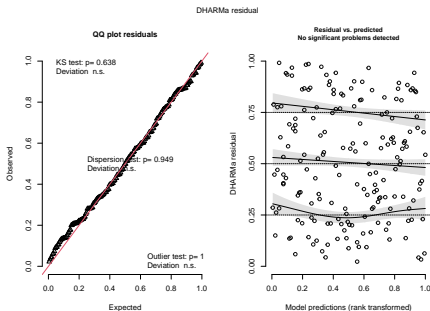
Quantile residuals

1. Simulate from posterior predictive distribution
 - ▶ `brms` has a `posterior_predict` function to generate samples in a matrix
2. Use `createDHARMA` function in the `DHARMA` package to format output
 - ▶ Input is a matrix of posterior samples and the observed outcome data
3. Make plots as before
 - ▶ Residuals vs predicted
 - ▶ Residuals vs predictors

DHARMa examples

```
postpred_sample_mod01 = posterior_aedaat %>% ungroup() %>%
  mutate(quantile_residual =
    ggplot(aes(x=BWT, y=quantile_residual)) +
    geom_point() +
    geom_smooth()
```

```
plot(dharma_resids)
```



Posterior predictive checks

1. Simulate from posterior predictive distribution
 - ▶ `tidy_bayes` has an `add_predicted_draws` function to append the samples to a data frame
2. Compute some summary statistic on each posterior draw and on the observed data
 - ▶ Summary statistic depends on what you want to diagnose
 - ▶ For binary models, it is almost always the expected value (mean)
3. Plot distribution of summary statistics and overlay observed values
 - ▶ Type of plot depends on grouping factor and summary statistic

Simulate from posterior predictive distribution

```
aedat_pp = add_predicted_draws(newdata = aedat, mod01_stan)

aedat_pp %>% ungroup() %>%
  select(USUBJID, PTTYPE, AE01, Quartile:.prediction) %>%
  slice_tail(n=4)
```

```
. # A tibble: 4 x 9
```


	USUBJID	PTTYPE	AE01	Quartile	.row	.chain	.iteration
	<fct>	<fct>	<int>	<chr>	<int>	<int>	<int>
. 1	UID-180	PT1	0	Q2	180	NA	NA
. 2	UID-180	PT1	0	Q2	180	NA	NA
. 3	UID-180	PT1	0	Q2	180	NA	NA
. 4	UID-180	PT1	0	Q2	180	NA	NA

Plot the VPC

```
# Observed data summary
obs_summary <- aedat %>%
  group_by(Quartile) %>%
  summarise(phat_obs = mean(AE

# Simulated data summary
sim_summary <- aedat_pp %>%
  group_by(.draw, Quartile) %>%
  summarise(phat_sim = mean(.p

# VPC
sim_summary %>%
  ggplot(aes(x=Quartile, y=pha
  geom_violin() +
  geom_point(data=obs_summary,
  labs(x='', y='Proportion wit
  coord_flip()
```



VPC for continuous variable: define summary statistic

- ▶ Fit generalized additive model (smoother) to each simulated dataset
- ▶ Predict at a fixed grid of values (0 to 95th percentile)

```
summary_function <- function(.data, .x_name, .y_name='value') {  
  .data <- .data %>% ungroup() %>% rename('xvar' = .x_name)  
  x_grid <- with(.data, seq(from = min(xvar),  
                           to = quantile(xvar, probs = 0.95),  
                           length = 100))  
  fit <- gam(yvar ~ s(xvar), family=binomial(link='logit'))  
  predictions <- predict(fit, newdata = data.frame(xvar=x_grid))  
  return( data.frame(xvar=x_grid, prediction = predictions))  
}
```

Compute summary statistics for observed data

```
obs_summary <- summary_function(aedat, .x_name = 'CAVGSS',  
  mutate(type='Observed'))
```

```
head(obs_summary)
```

.	xvar	prediction	type
1	0.00000000	0.03526032	Observed
2	0.03718346	0.03644987	Observed
3	0.07436693	0.03767799	Observed
4	0.11155039	0.03894582	Observed
5	0.14873385	0.04025454	Observed
6	0.18591731	0.04160537	Observed

Plot VPC

```
sim_summary %>% bind_rows(obs_summary) %>%  
  ggplot(aes(x=xvar, y=prediction)) +  
  geom_line(aes(col=type, group=type)) +  
  geom_ribbon(aes(ymin=qlo, ymax=qhi, fill=type), alpha=0.2)  
  labs(x='Steady-state Cavg', y='Probability of severe AE')
```



0.6

Model Comparison

- ▶ Goal: maximize expected log predictive density (ELPD) for future data
 - ▶ This is a measure of out-of-sample prediction quality
- ▶ Leave-one-out cross-validation (LOO-CV) to approximate ELPD
 - ▶ Involves fitting N models
 - ▶ Can be approximated using pareto smoothed importance sampling of the posterior samples
 - ▶ loo package
- ▶ WAIC also approximates -2 ELPD
 - ▶ Proven to be asymptotically equivalent to LOO-CV (modulo the -2)
 - ▶ Lower is better
 - ▶ LOO-CV generally preferable due to its ability to tell when the estimates are not trustworthy

Workbook Bayes02: Model evaluation and comparison

References